StandAId

RAG & PDF Extraction Specialist — Paid Trial Task

Fix table extraction & retrieval accuracy in our AI standards-search app

> **StandAId** is an app for Australian tradies: upload AS/NZS technical standards, AI extracts and chunks them for voice/text search, plus a learning & exam feature and onsite trade tools. We're looking for a specialist to harden the pipeline that turns dense, table-heavy PDF standards into accurate, searchable data.

Summary

We need a RAG/PDF-extraction specialist for a **paid trial task** first — get one real, scoped bug fixed well, then talk about ongoing work. This is precision extraction work: these documents contain safety-critical numbers (cable ratings, clearances, fault-loop impedance), so accuracy matters more than speed.

Our Stack

- Supabase (Postgres + pgvector, HNSW)
- Edge Functions — Deno / TypeScript
- OpenAI Embeddings
- Claude (Anthropic) — generation, vision, rerank
- Hybrid Retrieval — vector + full-text (RRF)

PDFs are scanned or digital, table-heavy, with clauses, appendices, and cross-references between standards.

Trial Task (pick one — scoped tight so it's fairly judgeable)

Option A — Table detection without a caption: reliably detect and transcribe tables that lack the literal word "TABLE" (headed only by a numbered clause title, e.g. "8.1 MAXIMUM VALUES OF...").

Option B — Missing-page recovery: detect PDF pages an OCR pass silently drops entirely (zero chunks land on that page), and build a retry/recovery path.

Option C — Retrieval eval harness: a repeatable script that scores retrieval accuracy against a small set of known Q&A pairs, so improvements can be measured, not guessed at.

Budget

$300 – $800 AUD for the trial task, fixed price — scoped exactly after a short call.

Ongoing work (hourly or per-task) to follow if the trial goes well.

Ideal Candidate

- Real production RAG experience — not just a LangChain demo
- Strong PDF parsing background, table extraction specifically (pdf.js / unpdf / pdfplumber, or vision-model OCR)
- **Hands-on LLM/AI API experience** — embeddings (OpenAI or similar), prompt design for extraction/reranking, vision models for OCR/table transcription (Claude or GPT-4V class)
- Comfortable with pgvector / Postgres full-text search and hybrid retrieval (RRF, reranking)
- TypeScript/Deno (Supabase edge functions), or fast to pick it up
- Explains tradeoffs in plain English — I'm an electrician by trade, not a software engineer, so clear communication matters as much as code quality

Screening Questions

"Describe a RAG system you've built where document structure (tables, columns, footnotes) actually mattered — what broke and how did you fix it?"

"How would you evaluate whether a retrieval change actually improved accuracy, rather than just 'felt better'?"